

Court Vision: Motion Planning for an Obstacle-Aware Ball-Dribbling Mobile Robot

RBE 550 Motion Planning — Project Proposal

Rudy Arsenault, Tracy Yang
Worcester Polytechnic Institute

Abstract—We propose a mobile robot that must dribble a ball with a paddle end-effector down a cluttered course, choosing where to strike the ball next so that the resulting bounce location is both reachable and collision-free. Unlike prior robot-dribbling work, which stabilizes a periodic bounce in place, or legged locomotion work, which dribbles across open terrain, our robot must jointly plan *where the ball goes* and *where the base goes* under hard obstacle constraints, occasionally redirecting the ball itself rather than only steering around obstacles. We outline the hybrid dribble dynamics, the coupled interception/navigation planner, a simulation-based validation plan in Gazebo or PyBullet, and a week-by-week schedule with contingencies scoped to a semester-length, sim-only team project.

I. Introduction

Ball dribbling is a deceptively rich motion planning problem. A player does not simply track the ball; the player decides, on every bounce, where to send it next, weighing the defender, the open lane, and where their own feet need to be a half-second later.

We propose to build a simulated mobile robot with a single paddle end-effector that must sustain a bouncing ball down a corridor cluttered with obstacles. The robot cannot simply avoid the clutter with its base, because the base’s position is coupled to the ball’s bounce location: if the nominal next-bounce spot is blocked, the robot must instead *redirect* the ball to a different, reachable, collision-free spot, and only then plan its base motion to intercept it there. This makes the problem a genuine hybrid of periodic contact control (the dribble) and spatial planning under constraints (the obstacles) with no single fixed goal configuration to plan toward. The “goal” is a moving target that the robot itself is allowed to choose within limits.

II. Background and Related Work

Robotic ball dribbling has been studied as a variant of robot juggling. Bühler, Koditschek, and Kindlmann formalized juggling and catching as intermittent dynamical tasks and introduced the “mirror law,” in which the end-effector’s motion is a distorted reflection of the ball’s motion so that a stable periodic bounce emerges from a simple feedback rule [1]. Bätz et al. showed that basketball-style dribbling can be modeled with the same hybrid (discrete impact / continuous flight) dynamics as juggling, and demonstrated stable multi-cycle dribbling on a stationary industrial arm using both force/torque and vision-based ball tracking [2, 3]. These works establish the dribble controller itself, but the manipulator base never moves and there is no notion of an obstacle to avoid.

Legged locomotion research has since pushed dribbling onto a moving platform: DribbleBot uses reinforcement learning to let a quadruped dribble a soccer ball across natural, uneven terrain (sand, mud, snow) while walking [4]. This demonstrates that dribbling and locomotion can be coupled, but the “obstacles” are terrain properties that perturb

the bounce, not discrete objects that must be routed around. There is no explicit path planning layer.

On the navigation side, classical approaches such as the dynamic window approach [5] and velocity obstacles [6] plan collision-free motion for a base among static and moving obstacles, but they assume the robot’s goal location is externally given and fixed at each planning cycle.

Our project sits in the gap between these three bodies of work: we adopt the hybrid dribble model of [1, 2], must move the base as in [4], but, unlike any of the above, must also treat the next-bounce target itself as a planning variable that can be shifted to satisfy obstacle constraints, coupling target selection with navigation-style path planning rather than treating them as sequential stages.

III. Proposed Approach

We will simulate a differential-drive base carrying a 1-DOF vertically actuated paddle in Gazebo or PyBullet. The ball is modeled as a sphere with a coefficient of restitution; between contacts it follows free-fall projectile motion, and at each paddle contact its outgoing velocity is set by the mirror-law-style controller of [1, 2].

Ball state is tracked in closed loop rather than predicted from a precomputed terrain model: after every bounce, an onboard camera estimate of the ball’s position and velocity is used to re-fit its projectile trajectory (in the spirit of the recursive tracking used in [3]), so that irregularities such as a tilted surface or an imperfectly known coefficient of restitution are absorbed by re-estimation rather than assumed away. Known obstacle geometry is used offline to build the course’s configuration-space map, but the ball’s own dynamics are always treated as sensed and re-estimated, not scripted.

The planner operates in two coupled layers, illustrated in Fig. 1:

1. **Bounce/interception layer:** given the ball’s current apex state, predict its default landing point and time, and check that point against the local obstacle map (built from the course’s known/perceived clutter, represented in configuration space as in the course’s C-space module).




Figure 1: The nominal next-bounce trajectory (dashed) is interrupted by an obstacle, so the redirect layer chooses a new, reachable landing point (solid arc) and the navigation layer re-plans the base pose to intercept it there.

2. **Redirect/navigation layer:** if the default landing point is blocked, search a small neighborhood of alternative landing points (by adjusting the paddle’s hit angle and impulse) for one that is both physically reachable by the dribble controller and collision-free; then plan a short-horizon, collision-free base trajectory to be under that point in time for the next contact.

This is the central creative and technical challenge of the project: most motion planning projects treat “avoid the obstacle” and “reach the goal” as the whole problem, but here the robot may also negotiate with its own target, which turns a single-shot planning query into a small per-cycle optimization repeated at dribble frequency.

IV. Objectives

1. **Dribble dynamics & controller:** implement the hybrid ball-paddle contact model and reproduce a stable, stationary periodic dribble in simulation.
2. **Base integration:** couple the controller’s commanded next contact point/time to a mobile-base trajectory so the robot can dribble while translating in an obstacle-free lane.
3. **Static obstacle course:** build a configuration-space map of a course with 4–6 static obstacles and plan collision-free base paths that sustain the dribble through the full course.
4. **Redirect-the-bounce planner:** when the nominal landing spot conflicts with an obstacle, compute an alternate hit angle/impulse that shifts the landing point to an open cell, and re-plan the base intercept accordingly, rather than only steering the base around obstacles.
5. **Evaluation:** quantify performance versus an obstacle-free baseline using metrics such as dribble cycles survived, path-length overhead, and redirect-planner success rate across obstacle densities.

V. Stretch Goals

1. **Dynamic/random clutter:** obstacles that appear or move during a run, requiring online re-planning rather than a single offline course map.
2. **Multi-bounce lookahead:** plan two to three bounces ahead, similar to how a real player reads the court, in-

stead of only the immediate next contact.

3. **Physical feasibility discussion:** document what would hypothetically need to change (sensing, paddle actuator bandwidth, ball tracking) to carry the approach onto real hardware.

VI. Schedule and Division of Labor

The schedule is aligned to the official RBE 550 course calendar with Week 8’s project status update and Week 16’s final report as hard checkpoints. Week 14 is deliberately reserved as a buffer; if the redirect-planner (Objective 4) is behind schedule at that point, the contingency plan is to fall back to a simpler policy in which the robot briefly catches and re-drops the ball to reset the bounce phase whenever the nominal target is blocked, rather than shifting the hit angle, which is a strictly easier version of the same idea that still produces a working, demonstrable system.

Wk	Objectives
3	Submit proposal; finalize team roles; stand up sim environment; deep-read dribbling/juggling and obstacle-avoidance literature.
4–5	Implement ball-paddle contact model and mirror-law-style controller; validate stationary periodic dribble (Obj. 1).
6–7	Attach base and couple contact point to base trajectory; build C-space map for a static obstacle course; first collision-free dribble-and-drive runs (Obj. 2–3).
8	Project status update: demonstrate baseline dribble + static-obstacle navigation.
9–10	Design and implement redirect-the-bounce planner; begin extending to moving obstacles (Obj. 4).
11–12	Robustness passes; failure-mode testing; metrics/logging pipeline (Obj. 5).
13	Full evaluation sweep across obstacle densities; begin stretch goals if ahead of schedule.
14	Buffer week / contingency fallback if needed (see above).
15	Finalize results and demo video; begin final report.
16	Final report and presentation due.

Division of labor: Arsenault will own the dribble dynamics, contact controller, and ball state estimation (Obj. 1, part of Obj. 4); Yang will own the configuration-space obstacle mapping, base path planning, and simulation/evaluation infrastructure (Obj. 2–3, Obj. 5). Both members will jointly design and implement the redirect-the-bounce planner (Obj. 4), since coupling the two halves of the system is the core intellectual contribution of the project, and will co-write the final report.

VII. Expected Results

We expect a simulation demo showing the robot sustaining a dribble through a static obstacle course that a purely base-avoidance planner (one that never adjusts the ball’s landing spot) cannot complete without dropping the ball or stalling. We will report, as a function of obstacle density, the dribble-

survival rate, the extra path length incurred by redirect events versus the obstacle-free baseline, and qualitative video comparisons. If stretch goals are reached, we will additionally report performance under randomly moving clutter and with multi-bounce lookahead.

References

- [1] M. Bühler, D. E. Koditschek, and P. J. Kindlmann, “Planning and control of robotic juggling and catching tasks,” *The International Journal of Robotics Research*, vol. 13, no. 2, pp. 101–118, 1994.
- [2] G. Bätz, M. Sobotka, D. Wollherr, and M. Buss, “Robot basketball: Ball dribbling – a modified juggling task,” in *Advances in Robotics Research*, T. Kröger and F. M. Wahl, Eds. Berlin, Heidelberg: Springer, 2009, pp. 323–334.
- [3] G. Bätz, K.-K. Lee, D. Wollherr, and M. Buss, “Robot basketball: A comparison of ball dribbling with visual and force/torque feedback,” in *Proc. IEEE Int. Conf. on Robotics and Automation (ICRA)*, 2009.
- [4] Y. Ji, G. B. Margolis, and P. Agrawal, “DribbleBot: Dynamic legged manipulation in the wild,” in *Proc. IEEE Int. Conf. on Robotics and Automation (ICRA)*, 2023.
- [5] D. Fox, W. Burgard, and S. Thrun, “The dynamic window approach to collision avoidance,” *IEEE Robotics & Automation Magazine*, vol. 4, no. 1, pp. 23–33, 1997.
- [6] P. Fiorini and Z. Shiller, “Motion planning in dynamic environments using velocity obstacles,” *The International Journal of Robotics Research*, vol. 17, no. 7, pp. 760–772, 1998.